

Linear Regression 1

Simple Linear Regression

Daniela Palleschi

2023-04-13

Table of contents

Resources	2
(Linear) Regression	2
Types of regression	3
Straight lines	3
Slopes	6
Slopes	6
Intercepts	6
Varying slopes and intercepts	6
A simple linear model	7
fit model	7
Coefficients table with <code>summary()</code>	7
Understanding the summary	8
.	10
Exploring the model	12
Assumptions	13
Normality assumption	14
Communicating your results	18

```
## play sound if error encountered
### from: https://sejohnston.com/2015/02/24/make-r-beep-when-r-markdown-finishes-or-when-i
options(error = function(){ # Beep on error
  beepr::beep(sound = "wilhelm")
  Sys.sleep(2) #
})
```

```

)
## and when knitting is complete
.Last <- function() {          # Beep on exiting session
  beeper::beep(sound = "ping")
  Sys.sleep(6) # allow to play for 6 seconds
}

# Create references.json file based on the citations in this script
# make sure you have 'bibliography: references.json' in the YAML
rbbt::bibt_update_bib("_lin_reg1.qmd")

```

Wrote 5 references to './references/references.json'

```

knitr::opts_chunk$set(eval = T, # change this to 'eval = T' to reproduce the analyses; make
                        echo = T, # 'print code chunk?'
                        message = F, # 'print messages (e.g., warnings)?'
                        error = F,
                        warning = F)

# suppress scientific notation
options(scipen=999)

# load libraries
library(tidyverse)
library(ggplot2)

```

Resources

- these slides are based on a mix of the following resources

DeBruine and Barr (2021); Winter (n.d.); Winter (2014); Winter (2019)

- and on slides that were originally based on Field, Miles, and Field (2013), but if you're looking for a textbook I'd recommend Winter (2019)

(Linear) Regression

- our data exploration has given us an idea about what our data look like

- we fit a model to our data, and use it to *predict* values of our DV based on one (or more) IV(s)
 - i.e., *predicting* an outcome variable (DIV) from one or more predictors (IVs)
- because we're making predictions, we need to take into account the variability (i.e., *error*) in our data

$$outcome_i = (model) + error_i$$

Types of regression

- Simple regression
 - single continuous predictor
- Multiple regression
 - multiple predictors
- Logistic regression
 - binary predictor
- Hierarchical/mixed models
 - include random effects

Straight lines

- regression: summarise the data with a straight line
- straight lines can be defined by
 - Slope (b_1)
 - * regression coefficient for the predictor
 - * gradient (slope) of the regression line
 - * direction/strength of relationship
 - Intercept (b_0)
 - * value of Y when $X = 0$

```
df_crit_verb <- readr::read_csv(here::here("data/tidy_data_lifetime_pilot.csv"),
                                # for special characters
                                locale = readr::locale(encoding = "latin1")
                                ) |>
```

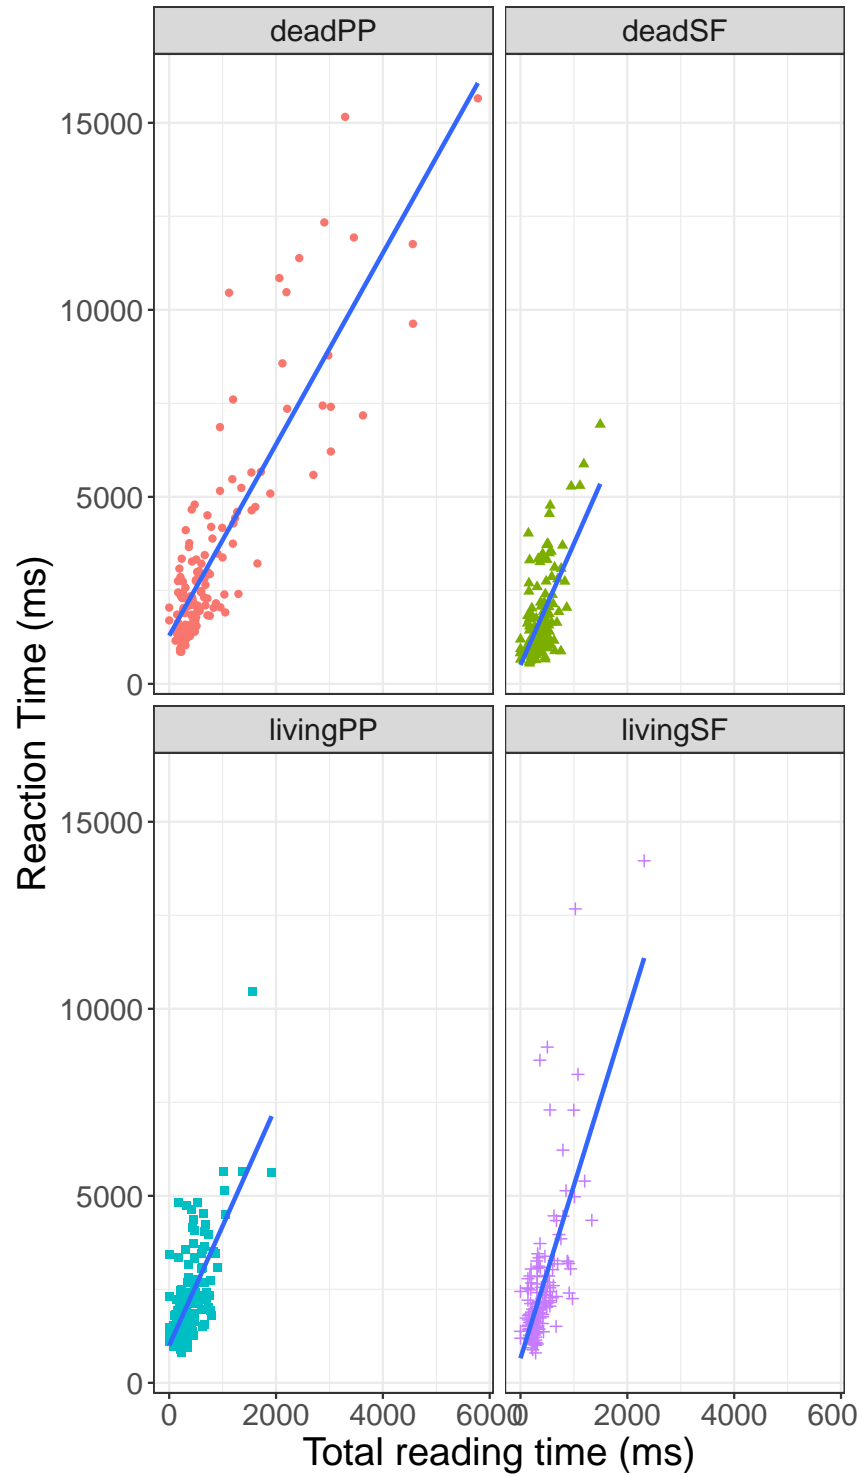
```

mutate_if(is.character, as.factor) |> # all character variables as factor
filter(type == "critical", # only critical trials
       px != "px7",
       region == "verb")

df_crit_verb |>
ggplot(aes(x = tt, y = rt)) +
  facet_wrap(~condition) +
  labs(title = "Scatterplots with regression line",
       x = "Total reading time (ms)",
       y = "Reaction Time (ms)") +
  geom_point(aes(colour = condition, shape = condition)) +
  geom_smooth(method="lm", se=F, fullrange=FALSE, level=0.95) +
  theme_bw() +
  theme(legend.position = "none",
       text = element_text(size=18))

```

Scatterplots with regression line



Slopes

- slopes describe a change in x (Δx) over a change in y (Δy)
 - positive slope: as x increases, y increases
 - negative slope: as x increases, y decreases
 - if the slope is 0, there is no change in y as a function of x

$$\text{slope} = \frac{\Delta x}{\Delta y}$$

- or: the change in y when x increase by 1 unit
 - sometimes referred to as “rise over run”: how do you ‘rise’ in y for a given ‘run’ in x ?

Slopes

Intercepts

Varying slopes and intercepts




Figure 1: Image source: Winter (2019) (all rights reserved)

A simple linear model

- first, order our predictor
 - we predict longer reading times for dead versus living, so order living dead

```
# order factor levels
df_crit_verb$lifetime <- factor(df_crit_verb$lifetime, levels = c("living","dead"))

# set contrasts
contrasts(df_crit_verb$lifetime) <- c(-0.5,+0.5); contrasts(df_crit_verb$lifetime)
```

```
      [,1]
living -0.5
dead    0.5
```

fit model

- let's exclude missing observations (0)

```
# fit simple linear model
fit_1 <- df_crit_verb %>%
  filter(ff > 0) %>%
  lm(ff ~ lifetime, data = .)

# alternatively
fit_1 <- lm(ff ~ lifetime,
            data = df_crit_verb, subset = ff > 0)
```

Coefficients table with summary()

```
> summary(fit_1)

Call:
lm(formula = ff ~ lifetime, data = df_crit_verb, subset = ff > 0) #<1>

Residuals:                                     #<2>
    Min       1Q   Median       3Q      Max
-118.78  -37.78  -11.39   25.22  264.61
```

```

Coefficients:
              Estimate Std. Error t value Pr(>|t|)      #<3>
(Intercept)   199.089      2.466  80.743  <2e-16 ***    #<4>
lifetimedeath  5.388       4.931   1.093   0.275          #<5>
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 57.46 on 541 degrees of freedom
Multiple R-squared:  0.002202, Adjusted R-squared:  0.0003575    #<6>
F-statistic: 1.194 on 1 and 541 DF,  p-value: 0.275

```

- ① formula
- ② Residuals: differences between observed values and those predicted by the model
- ③ Names for columns Estimates, SE, t-value, p-value
- ④ Intercept (b_0), i.e., value of y (first fix.) with a move of one unit of x (lifetime)
- ⑤ Slope (b_1), i.e., change in first fixation going from **dead** to **living**
- ⑥ (Adjusted) R^2 , a measure of model fit (squared residuals)
 - what is the **intercept**?
 - is the **slope** positive or negative?
 - what is it's value?
 - this is what the slope would look like:

Understanding the summary

- let's compute summary statistics based on *lifetime*
 - then compare this to the model output

Exercises

1. Subtract the mean first-fixation reading time of **dead** from that of **living**
 - what does this correspond to in the model summary?
2. Compute the mean of **dead+living**
 - what does this correspond to in the model summary?
3. Divide the slope in 2. Subtract this from the mean of **dead**.
 - what does this correspond to?

Summary statistics

Table 1: Summmary statistics for first-fixation duration at the verb region

lifetime	N	mean	sd	se	ci	lower.ci	upper.ci
living	271	196.395	59.639	3.623	7.133	189.262	203.528
dead	272	201.783	55.197	3.347	6.589	195.194	208.372

```
# compute summary
summary_ff_life <- df_crit_verb |>
  filter(region=="verb",
         ff > 0) |>
  group_by(lifetime) %>%
  summarise(N = n(),
            mean = mean(ff, na.rm = T),
            sd = sd(ff, na.rm = T)) %>%

# compute standard error, confidence intervals, and lower/upper ci bounds
mutate(se = sd / sqrt(N),
       ci = qt(1 - (0.05 / 2), N - 1) * se,
       lower.ci = mean - qt(1 - (0.05 / 2), N - 1) * se,
       upper.ci = mean + qt(1 - (0.05 / 2), N - 1) * se)

knitr::kable(summary_ff_life, digits=3,
              caption = "Summmary statistics for first-fixation duration at the verb region",
              kableExtra::kable_styling(font_size = 24,
                                         position = "left"))
```

Model summary

```
summary(fit_1)
```

Call:

```
lm(formula = ff ~ lifetime, data = df_crit_verb, subset = ff >
0)
```

Residuals:

Min	1Q	Median	3Q	Max
-118.78	-37.78	-11.39	25.22	264.61

Coefficients:

Estimate	Std. Error	t value	Pr(> t)

```
(Intercept) 199.089      2.466  80.743 <0.0000000000000002 ***
lifetime1    5.388      4.931   1.093              0.275
```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 57.46 on 541 degrees of freedom

Multiple R-squared: 0.002202, Adjusted R-squared: 0.0003575

F-statistic: 1.194 on 1 and 541 DF, p-value: 0.275

💡 broom package for tidy model summaries

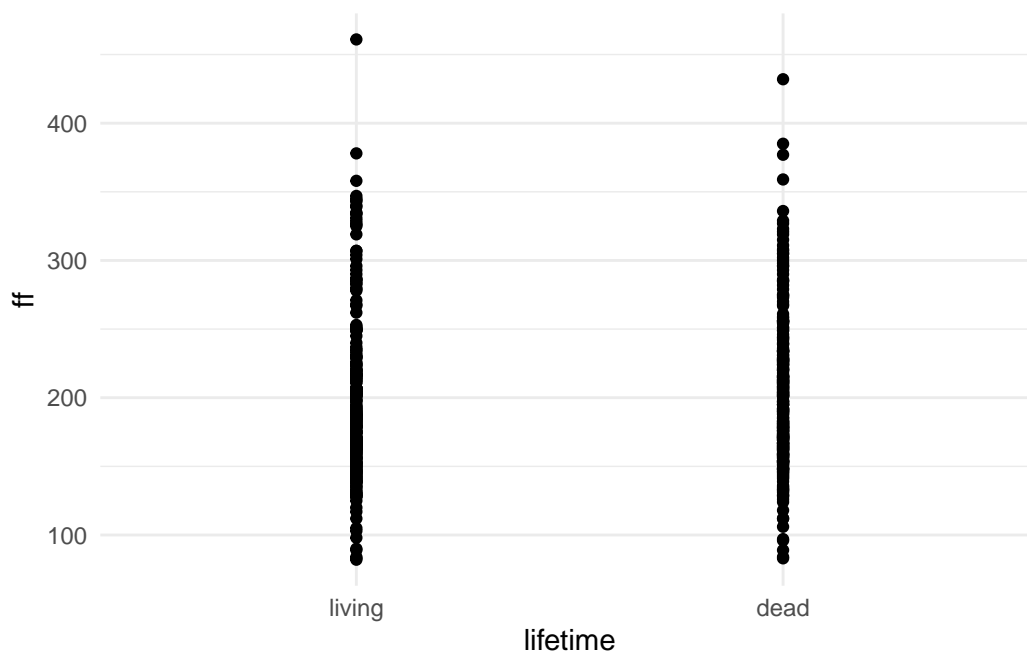
```
# install.packages("broom")
library(broom)
tidy(fit_1)

# A tibble: 2 x 5
  term          estimate std.error statistic    p.value
<chr>         <dbl>     <dbl>     <dbl>    <dbl>
1 (Intercept)  199.        2.47      80.7 5.95e-304
2 lifetime1     5.39      4.93       1.09 2.75e- 1

glance(fit_1)

# A tibble: 1 x 12
  r.squared adj.r.squared sigma statistic p.value    df logLik   AIC   BIC
  <dbl>      <dbl>    <dbl>     <dbl>   <dbl> <dbl> <dbl> <dbl> <dbl>
1  0.00220    0.000358  57.5        1.19  0.275     1 -2969. 5944. 5957.
# i 3 more variables: deviance <dbl>, df.residual <int>, nobs <int>

df_crit_verb %>%
  filter(ff > 0) |>
  ggplot(aes(x = lifetime, y = ff)) +
  geom_point(position = position_dodge(.6)) +
  geom_smooth(method = 'lm') + theme_minimal()
```



```
df_crit_verb %>%
  # mutate(life_c = if_else(lifetime=="living",-0.5,+0.5)) %>%
  select(ff,fp,rpd,tt) %>%
  cor()
```

	ff	fp	rpd	tt
ff	1.0000000	0.4530152	0.4146744	0.1273391
fp	0.4530152	1.0000000	0.8956934	0.3620559
rpd	0.4146744	0.8956934	1.0000000	0.3325920
tt	0.1273391	0.3620559	0.3325920	1.0000000

```
# fit simple linear model with log
fit_2 <- df_crit_verb %>%
  filter(ff > 0) %>% # because you can't log transform 0
  mutate(ff_c = ff-mean(ff)) %>%
  lm(ff_c ~ lifetime, data = .)
summary(fit_2)
```

Call:
lm(formula = ff_c ~ lifetime, data = .)

```

Residuals:
      Min       1Q   Median       3Q      Max
-118.78  -37.78  -11.39   25.22  264.61

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) -0.004962    2.465724  -0.002   0.998
lifetime1    5.388254    4.931448   1.093   0.275

Residual standard error: 57.46 on 541 degrees of freedom
Multiple R-squared:  0.002202, Adjusted R-squared:  0.0003575
F-statistic: 1.194 on 1 and 541 DF, p-value: 0.275

```

Exploring the model

- the linear model contains *fitted* values corresponding to our *observed* values
 - these *fitted* values are fit to a straight line
 - our *observed* values are not fit to a straight line
 - the *residuals* are the differences along the *y* axis from the fitted to the observed values

```

# how many observed values did we enter into the model?
df_crit_verb |>
  filter(ff > 0) |>
  nrow()

```

[1] 543

```

# how many observed values did we enter into the model?
length(fitted(fit_1))

```

[1] 543

```

# how many observed values did we enter into the model?
head(fitted(fit_1))

```

```

      1      2      3      4      5      6
196.3948 196.3948 201.7831 196.3948 201.7831 196.3948

```

```
# look at the residuals
head(residuals(fit_1))
```

1	2	3	4	5	6
-21.39483	10.60517	26.21691	34.60517	10.21691	-28.39483

```
coef(fit_1)
```

(Intercept)	lifetime1
199.088961	5.388254

```
coef(fit_1)['(Intercept)'] + coef(fit_1)['lifetime1'] * -0.5
```

(Intercept)
196.3948

```
coef(fit_1)['(Intercept)'] + coef(fit_1)['lifetime1'] * 0.5
```


(Intercept)
201.7831

Assumptions


- normality assumption
 - residuals of the model are (approximately) normally distributed
- constant variance assumption (homoscedasticity)
 - spread of residuals should be (approximately) equal along the regression line

Normality assumption

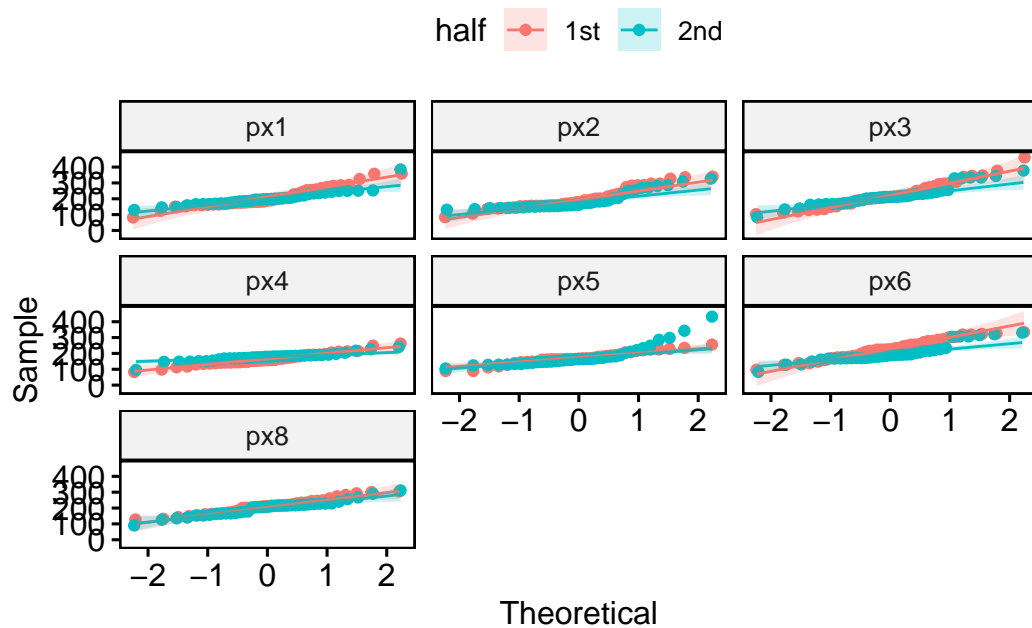
```
df_crit_verb |>  
  filter(ff > 0) |>  
  mutate(half = if_else(trial >= 104, "1st","2nd")) |>  
  ggpubr::ggqqplot(x = "ff")
```




```
df_crit_verb |>  
  filter(ff > 0) |>  
  mutate(half = if_else(trial >= 104, "1st","2nd")) |>  
  ggpubr::ggqqplot( x = "ff",  
                    color = "half")
```



```
df_crit_verb |>
  filter(ff > 0) |>
  mutate(half = if_else(trial >= 104, "1st", "2nd")) |>
  ggpubr::ggqqplot( x = "ff",
                    color = "half",
                    facet.by = "px")
```



```
plot(density(resid(fit_1)))
```



Log transformation

- for more see Section 5.4 in Winter (2019)
- the R function `log()` computes the ‘natural logarithm’ (and is the inverse of the exponential `exp()`)
 - `log()` makes large numbers smaller
 - `exp()` makes small numbers larger

```
log(0)
```

```
[1] -Inf
```

```
log(1:10)
```

```
[1] 0.0000000 0.6931472 1.0986123 1.3862944 1.6094379 1.7917595 1.9459101  
[8] 2.0794415 2.1972246 2.3025851
```

```
log(c(10,20,30,40,100))
```

```
[1] 2.302585 2.995732 3.401197 3.688879 4.605170
```

```
exp(1:10)
```

```
[1] 2.718282 7.389056 20.085537 54.598150 148.413159  
[6] 403.428793 1096.633158 2980.957987 8103.083928 22026.465795
```

```
exp(log(1:10))
```

```
[1] 1 2 3 4 5 6 7 8 9 10
```

- continuous variables truncated at 0 typically have a *positive skew*
 - a lot of small values (e.g., `tt` < 500ms), with some larger values (> `tt` 1000)
 - this usually means our residuals are also positively skewed, i.e., not normally distributed
- so we typically log-transform raw reading/reaction times for our linear models

```
# fit simple linear model with log
fit_2 <- df_crit_verb %>%
  filter(ff > 0) %>% # important! you can't log transform 0
  lm(log(ff) ~ lifetime, data = .)
summary(fit_2)
```

Call:

```
lm(formula = log(ff) ~ lifetime, data = .)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.85242	-0.17140	-0.01899	0.15691	0.89550

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	5.25458	0.01197	439.064	<0.0000000000000002 ***
lifetime1	0.03336	0.02394	1.394	0.164

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.2789 on 541 degrees of freedom


Multiple R-squared: 0.003579, Adjusted R-squared: 0.001737

F-statistic: 1.943 on 1 and 541 DF, p-value: 0.1639


Communicating your results

- model summaries can be provided via tables and/or figures
 - you should always report the t-values and p-values of an effect

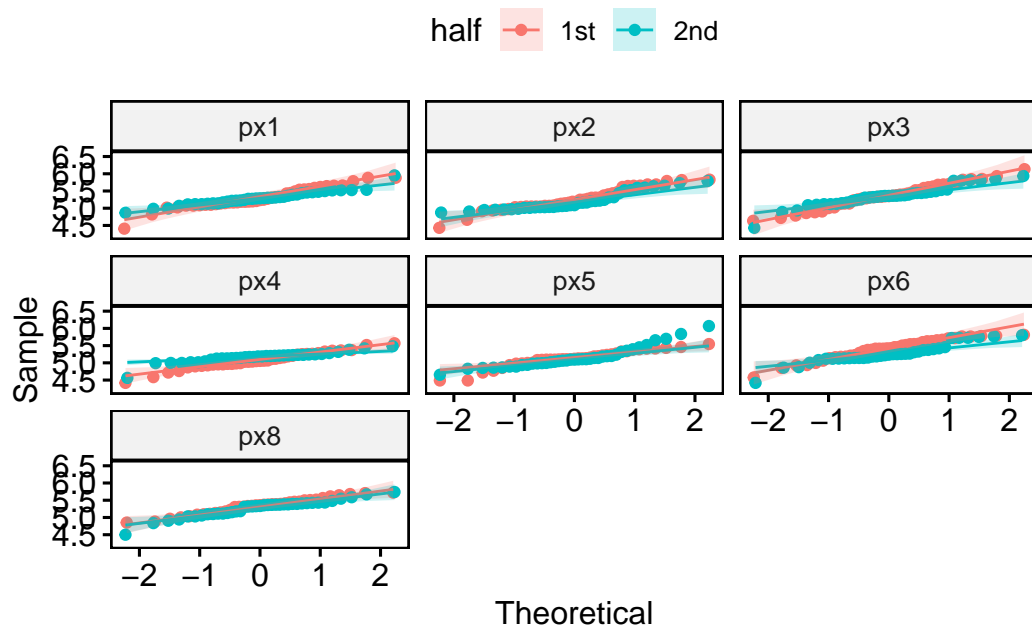
```
df_crit_verb |>
  filter(ff > 0) |>
  mutate(log_ff = log(ff)) |>
  mutate(half = if_else(trial >= 104, "1st", "2nd")) |>
  ggpubr::ggqqplot(x = "log_ff")
```




```
df_crit_verb |>
  filter(ff > 0) |>
  mutate(log_ff = log(ff)) |>
  mutate(half = if_else(trial >= 104, "1st", "2nd")) |>
  ggpubr::ggqqplot( x = "log_ff",
                    color = "half")
```



```
df_crit_verb |>
  filter(ff > 0) |>
  mutate(log_ff = log(ff)) |>
  mutate(half = if_else(trial >= 104, "1st", "2nd")) |>
  ggpubr::ggqqplot( x = "log_ff",
                    color = "half",
                    facet.by = "px")
```



```
plot(density(resid(fit_1)))
```



```

b0 <- coef(fit_2)["(Intercept)"]
b1 <- (coef(fit_2)["lifetime1"])

exp(b0+b1)

```

```

(Intercept)
197.9364

```

- DeBruine, Lisa M., and Dale J. Barr. 2021. "Understanding Mixed-Effects Models Through Data Simulation." *Advances in Methods and Practices in Psychological Science* 4 (1): 251524592096511. <https://doi.org/10.1177/2515245920965119>.
- Field, Andy, Jeremy Miles, and Zoe Field. 2013. *Discovering Statistics Using R*. Vol. 50. <https://doi.org/10.5860/choice.50-2114>.
- Winter, Bodo. 2014. "A Very Basic Tutorial for Performing Linear Mixed Effects Analyses (Tutorial 2)." 2014.
- . 2019. *Statistics for Linguists: An Introduction Using R*. Routledge. <https://doi.org/10.4324/9781315165547>.
- . n.d. "Linear Models and Linear Mixed Effects Models in R: Tutorial 1."